

✓ Congratulations! You passed!

Go to next item

Grade received 100% Latest Submission Grade 100% To pass 80% or higher

1. True/False: Suppose you learn a word embedding for a vocabulary of 60000 words. Then the embedding vectors could be 60000 dimensional, so as to capture the full range of variation and meaning in those words.

1 / 1 point

- ☐ True
- ☒ False

↗ Expand

✓ Correct

No, the dimension of word vectors is usually smaller than the size of the vocabulary. Most common sizes for word vectors range between 50 and 1000.

2. What is t-SNE?

1 / 1 point

- ☐ A linear transformation that allows us to solve analogies on word vectors
- ☐ An open-source sequence modeling library
- ☐ A supervised learning algorithm for learning word embeddings
- ☒ A non-linear dimensionality reduction technique

↗ Expand

✓ Correct

Yes

3. Suppose you download a pre-trained word embedding which has been trained on a huge corpus of text. You then use this word embedding to train an RNN for a language task of recognizing if someone is happy from a short snippet of text, using a small training set.

1 / 1 point

x (input text)	y (happy?)
I'm feeling wonderful today!	1
I'm bummed that my cat is ill.	0
Really enjoying this!	1

True/False: Then even if the word “upset” does not appear in your small training set, your RNN might reasonably be expected to recognize “I’m upset” as deserving a label $y = 0$.

- ☐ False
- ☒ True

↗ Expand

✓ Correct

Yes, word vectors empower your model with an incredible ability to generalize. The vector for “upset” would contain a negative/unhappy connotation which will probably make your model classify the sentence as a “0”.

4. Which of these equations do you think should hold for a good word embedding? (Check all that apply)

1 / 1 point

☐ $e_{boy} - e_{brother} \approx e_{sister} - e_{girl}$

☐ $e_{boy} - e_{girl} \approx e_{sister} - e_{brother}$

☒ $e_{boy} - e_{girl} \approx e_{brother} - e_{sister}$

✓ Correct
Yes!

☒ $e_{boy} - e_{brother} \approx e_{girl} - e_{sister}$

✓ Correct
Yes!

↗ Expand

✓ Correct
Great, you got all the right answers.

5. Let A be an embedding matrix, and let o_{4567} be a one-hot vector corresponding to word 4567. Then to get the embedding of word 4567, why don't we call $A * o_{4567}$ in Python?

1 / 1 point

- ☐ The correct formula is $A^T * o_{4567}$
- ☐ None of the answers are correct: calling the Python snippet as described above is fine.
- ☒ It is computationally wasteful.
- ☐ This doesn't handle unknown words (<UNK>).

↗ Expand

✓ Correct
Yes, the element-wise multiplication will be extremely inefficient.

6. When learning word embeddings, we create an artificial task of estimating $P(\text{target} \mid \text{context})$. It is okay if we do poorly on this artificial prediction task; the more important by-product of this task is that we learn a useful set of word embeddings.

1 / 1 point

- ☐ False
- ☒ True

↗ Expand

✓ Correct

7. True/False: In the word2vec algorithm, you estimate $P(t \mid c)$, where t is the target word and c is a context word. t and c are chosen from the training set to be nearby words.

1 / 1 point

- ☒ True
- ☐ False

Expand

Correct

Yes, and are chosen from the training set to be nearby words.

8. Suppose you have a 10000 word vocabulary, and are learning 500-dimensional word embeddings. The word2vec model uses the following softmax function:

1 / 1 point

$$P(t | c) = \frac{e^{\theta_t^T e_c}}{\sum_{t=1}^{10000} e^{\theta_t^T e_c}}$$

Which of these statements are correct? Check all that apply.

☒ θ_t and e_c are both 500 dimensional vectors.

Correct

☒ θ_t and e_c are both trained with an optimization algorithm such as Adam or gradient descent.

Correct

☐ After training, we should expect θ_t to be very close to e_c when t and c are the same word.

☐ θ_t and e_c are both 10000 dimensional vectors.

Expand

Correct

Great, you got all the right answers.

9. Suppose you have a 10000 word vocabulary, and are learning 500-dimensional word embeddings. The GloVe model minimizes this objective:

1 / 1 point

$$\min \sum_{i=1}^{10,000} \sum_{j=1}^{10,000} f(X_{ij})(\theta_i^T e_j + b_i + b_j - \log X_{ij})^2$$

True/False: X_{ij} is the number of times word j appears in the context of word i .

☒ True

☐ False

Expand

Correct

X_{ij} is the number of times word j appears in the context of word i .

10. You have trained word embeddings using a text dataset of t_1 words. You are considering using these word embeddings for a language task, for which you have a separate labeled dataset of t_2 words. Keeping in mind that using word embeddings is a form of transfer learning, under which of these circumstances would you expect the word embeddings to be helpful?

1 / 1 point

☒ When t_1 is larger than t_2

☐ When t_1 is smaller than t_2

☐ When t_1 is equal to t_2

Expand

Correct

Transfer embeddings to new tasks with smaller training sets.

